

Algorithms is “Cool”

- ❑ Where do we use Algorithms
- ❑ Analysis of Algorithms & Why it is important
- ❑ Sample Computational Problems

Why Algorithm is Cool:

Algorithms is Anywhere & Everywhere.

Motivation (Why algorithms?)

□ Where do we use Algorithms?

- ❖ Computer Science, Engineering
- ❖ Business, Operations Research
- ❖ Finance, Social Sciences, ... , Everywhere

□ Why is Performance Important?

- ❖ Exponential-size solution space
- ❖ NP-completeness

□ Problem Size Explosion

- ❖ Computer Chip Complexity,
- ❖ Database Sizes,
- ❖ Human Genome Project, WWW (Google),

Diverse Applications (where?)

- ❖ **Design the next generation CPU/GPU chip**
 - ◆ *how to design optimally (w.r.t. speed, power, area)*
- ❖ **Internet (WWW)**
 - ◆ *how to manage, manipulate large volume of data*
- ❖ **e-Commerce**
 - ◆ *how to manage transaction (secure, private)*
- ❖ **Logistics**
 - ◆ *how to manage transport/transfer of goods, people*
- ❖ **Human Genome Project**
 - ◆ *how to analyze huge volume of DNA/protein data*

Analysis of algorithms

The theoretical study of program performance and resource usage.

What's are the important aspects of software?

- modularity
- correctness
- maintainability
- functionality
- robustness
- user-friendliness
- programmer time
- simplicity
- extensibility
- reliability

** Speed / performance*

Why study algorithms and performance?

- ❑ Algorithms help us to understand *scalability*.
- ❑ Performance often draws the line between what is *feasible* and what is *impossible*.
- ❑ Performance is the *currency* of computing.
- ❑ The lessons of program performance *generalize* to other computing resources.
- ❑ Speed is *fun*!

Some Combinatorial Problems

Computational Complexity

Combinatorial Problem

| | |
|-------------------------------------|-------------------|
| ❖ Maintaining student records | <i>Easy</i> |
| ❖ Data Compression | <i>Easy</i> |
| ❖ Traveling Salesman Problem | <i>Hard</i> |
| ❖ Shortest Route Planning | <i>Easy</i> |
| ❖ Program Halting Problem | <i>Impossible</i> |
| ❖ VLSI Chip Layout Problem | <i>Hard</i> |
| ❖ Examination Time Table Scheduling | <i>Hard</i> |
| ❖ Checking if graph is acyclic | <i>Easy</i> |
| ❖ Computer Deadlock Problem | <i>Easy</i> |
| ❖ Sorting records in a Database | <i>Easy</i> |
| ❖ Finding patterns in a text | <i>Easy</i> |

Combinatorial Problems...

□ General Combinatorial Problem

- ❖ Given a finite, discrete set S of objects
- ❖ *To compute some function $f(S)$*

□ Algorithmic Issues...

- ❖ Representation of the set S
- ❖ Efficient manipulation of the set S
- ❖ Efficient algorithm to compute $f(S)$

Design and Analysis of Algorithms

□ Given a problem P ,

❖ Can it be solved?

Computability

□ If “Yes”, given an algorithm A for solving P ,

❖ Is algorithm A correct ?

Verification

❖ How good is algorithm A ?

Efficiency

❖ Can find a better algorithm A' ?

□ How do we define good?

❖ How much time it takes.

Time Complexity

❖ How much space it uses.

Space Complexity

Complexity

□ Given an algorithm A for problem P ,

□ How to do better?

Complexity of Algorithm

- ❖ Is the time complexity of A polynomial ?
- ❖ Can we design faster algorithm A' ?
- ❖ Can we design algorithm A'' that uses less space

□ Can we do better?

Complexity of Problem

- ❖ Is the problem NP-complete?
- ❖ Can we establish lower bounds
- ❖ Is the algorithm “*best possible*”

Why study program performance?



***But, is it useful
at all?***

**Q: Can we trust the CPUs in
our laptop / iPad / Apple Watch?**

Is the $MULT(*)$ operation correct?

Your laptop / iPad / Apple Watch
all have a CPU inside.

The CPU has a $MULT$ operation $(*)$



1994, Pentium FDIV bug

Pentium FDIV bug

From Wikipedia, the free encyclopedia

The **Pentium FDIV bug** was a **computer bug** that affected the **floating point unit** (FPU) of the **early Intel Pentium processors**. Because of the bug, the processor could return incorrect binary floating point results when dividing a number. Discovered in 1994 by Professor Thomas R. Nicely at **Lynchburg College**,^[1] Intel attributed the error to missing entries in the lookup table used by the floating-point division circuitry.^[2]

The severity of the FDIV bug is debated. Intel, producer of the affected chip, claims that the common user would experience it once every 27,000 years while IBM, manufacturer of a chip competing with Intel's Pentium, claims that the common user would experience it once every 24 days. Though rarely encountered by most users (*Byte* magazine estimated that 1 in 9 billion floating point divides with random parameters would produce inaccurate results),^[3]

both the flaw and Intel's initial handling of the matter were heavily criticized by the tech community. The man who



https://en.wikipedia.org/wiki/Pentium_FDIV_bug

What if we don't trust the CPU?



*How can we
check it...*

**Q: Can we trust the CPUs in
our laptop / iPad / Apple Watch?**

Testing the * operation in a CPU

Q: How to test that the “*” operation of your CPU is correct?

A: Check exhaustively. For all a, b
Check $a * b = c$

Q: How long will it take?

A: Any guesses?



Testing the * operation in a CPU

Q: How long will it take?

...very fast.
Laptop ~3G-Flop

Assume we use a 100G-Flop CPU
can do 100B operations per sec.

a is a 32-bit number (2^{32} cases)

b is a 32-bit number (2^{32} cases)

So, $(a * b)$ there are (2^{64} cases)

Time taken = $(2^{64} / 100 \times 10^9)$ sec





$2^{64} / 100 \times 10^9$ seconds



Examples Random

Assuming seconds of time for "seconds" | Use seconds of arc Instead

<http://www.wolframalpha.com/input/?i=2^64+%2F+100x10^9+seconds> ←- URL

Input interpretation:

$$\frac{2^{64}}{100 \times 10^9} \text{ seconds}$$

Unit conversions:

More

3.074×10^6 minutes

51 241 hours

2135 days

305 weeks

70.19 months

Testing $*$ operation in a CPU

Q: How long will it take?

Assume we use a 100G-Flop machine
can take 100B operations per sec.

a is a 32-bit number (2^{32} cases)

b is a 32-bit number (2^{32} cases)

So, $(a * b)$ there are (2^{64} cases)

Time taken = $(2^{64} / 100 \times 10^9)$ sec

≈ 6 years!



Summary

Testing * operation in a CPU

Q: If we use $O(n^2)$ algorithm?

≈ 6 years!



*Impossible.
Not practical*



Testing * operation in a CPU

Q: What if we have $(n \lg n)$ algorithm?

Assume we use a 100G-Flop machine
can do 100B operations per sec.

a is a 32-bit number (2^{32} cases)

b is a 32-bit number (2^{32} cases)

When $n = 2^{32}$, with $(n \lg n)$ algorithm

Time taken = $(2^{32} * 32 / 100 \times 10^9)$ sec

< 2 sec!



Summary

Testing * operation in a CPU

Q: If we use $O(n^2)$ algorithm?

≈ 6 years!



*Impossible.
Not practical*

Q: If we use $O(n \lg n)$ algorithm?

< 2 sec!



Fast & Practical



Moral of the story

Analysis of algorithms
help us make *predictions*.

Analysis of algorithms
help us *prepare for the worst case*.

Application in Web-Service

Suppose you code up a new web-service – *CoolApp*

- you debugged your code, and after some time.
- you got it *working*, you *tested it* a little bit
- it is quite fast

Can you release *CoolApp*?
Will it work? Or will it bomb?

If dream come true & *CoolApp*'s wildly popular?
How fast is “quite fast”, will server die? When?

Moral of the story

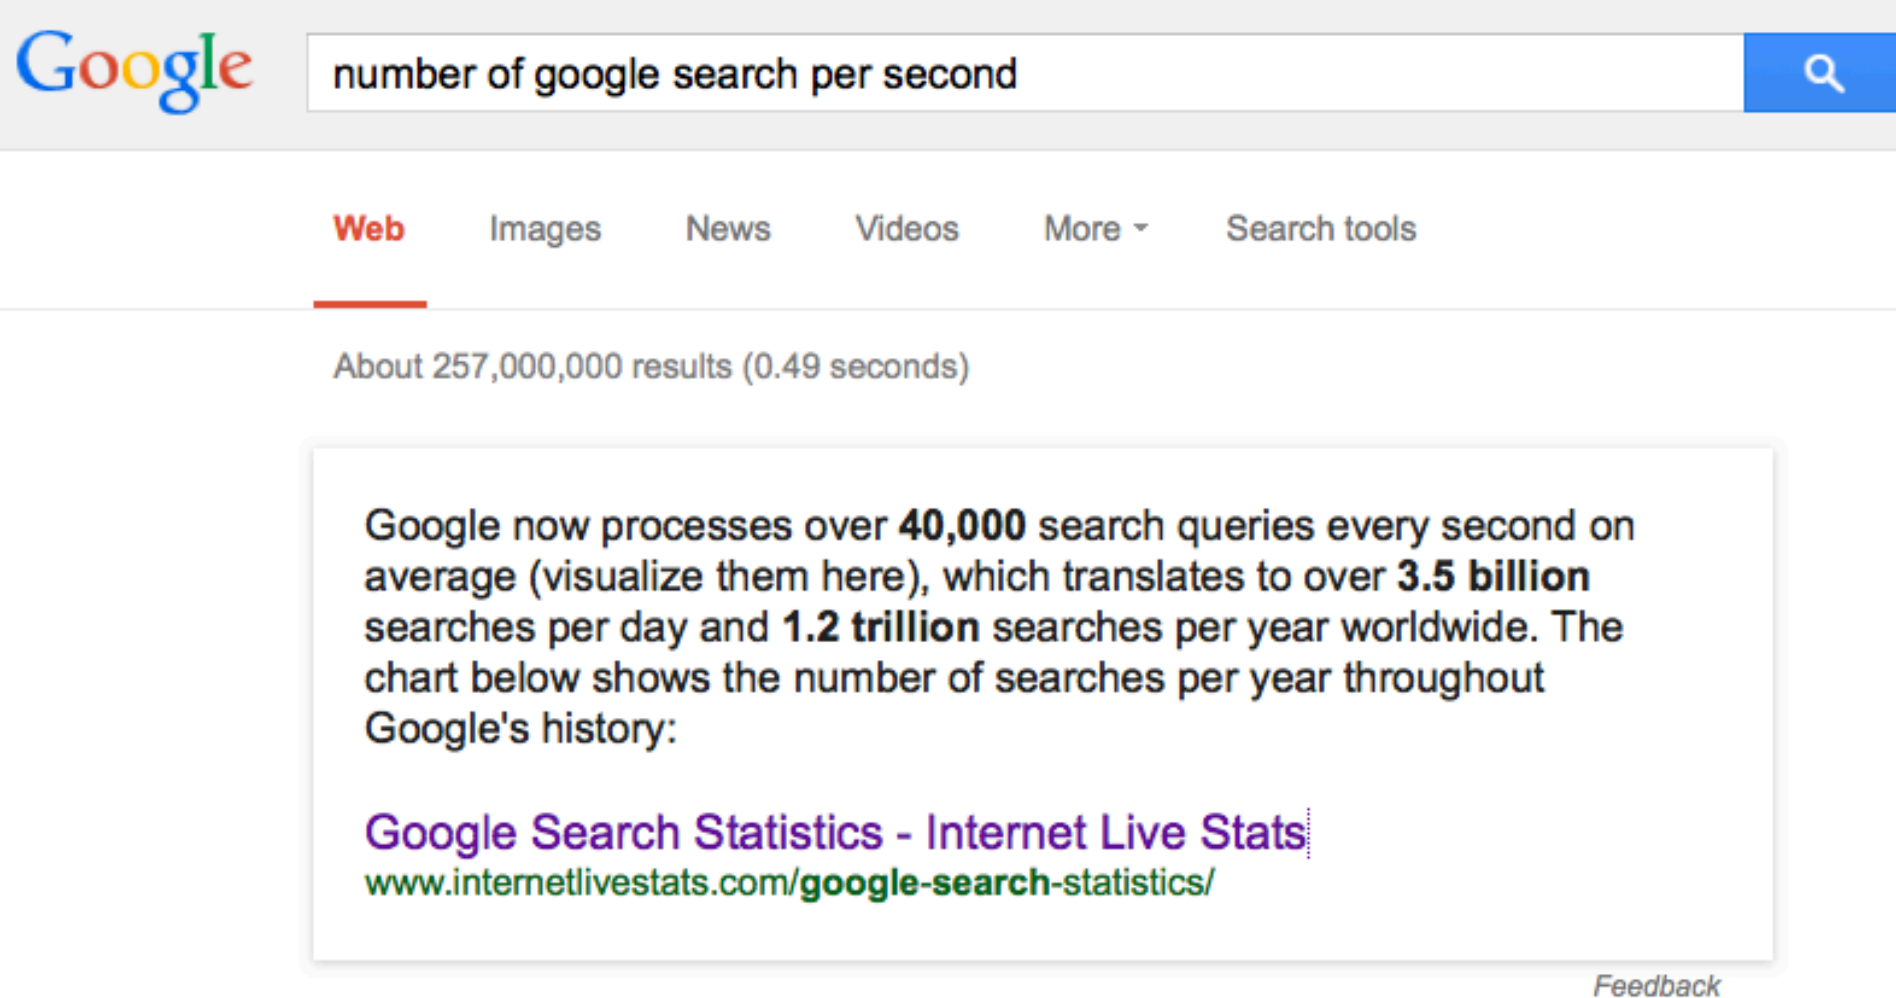
Analysis of algorithms
help us make *predictions*.

Analysis of algorithms
help us *prepare for the worst case*.

Note: If operation is “quite fast”, 0.02sec/op
that's 3min for 10,000 clicks per second,
that's 12min for 40,000 clicks per second,

Also, how big a load can a server take before dying?

40,000 clicks per seconds (July 2015)



The image is a screenshot of a Google search page. At the top left is the Google logo. To its right is a search bar containing the text "number of google search per second". A blue search button with a magnifying glass icon is to the right of the search bar. Below the search bar are navigation links: "Web" (highlighted with a red underline), "Images", "News", "Videos", "More ▾", and "Search tools". Below these links, it says "About 257,000,000 results (0.49 seconds)". The main content area contains a text block stating: "Google now processes over **40,000** search queries every second on average (visualize them here), which translates to over **3.5 billion** searches per day and **1.2 trillion** searches per year worldwide. The chart below shows the number of searches per year throughout Google's history:". Below this text is a link in purple: "Google Search Statistics - Internet Live Stats" followed by the URL "www.internetlivestats.com/google-search-statistics/" in green. At the bottom right of the content area is a "Feedback" link.

Google

number of google search per second

Web Images News Videos More ▾ Search tools

About 257,000,000 results (0.49 seconds)

Google now processes over **40,000** search queries every second on average (visualize them here), which translates to over **3.5 billion** searches per day and **1.2 trillion** searches per year worldwide. The chart below shows the number of searches per year throughout Google's history:

Google Search Statistics - Internet Live Stats
www.internetlivestats.com/google-search-statistics/

Feedback

Story of Algorithms in Action

Credit card processing centre in SG (Sci Park):

- monitor showing servers load for diff. countries,
 - blue, green, yellow, orange (send SMS alert),
- RED (URGENT Alert! → deploy more servers!)




Note: *Picture is NOT the real thing.*

But it gives the rough idea
and “demos” my point.

Note to Self:

Picture is NOT very good.
Will find a better one.



Why study algorithms and performance?

- Algorithms help us to understand *scalability*.
- Performance often draws the line between what is feasible and what is impossible.
- Algorithmic mathematics provides a *language* for talking about program behavior.
- Performance is the *currency* of computing.
- The lessons of program performance generalize to other computing resources.
- Speed is fun!

Thank you.

Q & A



School *of* Computing